

Operating Systems

Process Synchronization

Session 5

A Biswas

Process Synchronization

Process Synchronization

- **The Critical-Section Problem**
- **Peterson's Solution**
- **Bakery Algorithm**
- **Synchronization Hardware**
- **Semaphores**
- **Classic Problems of synchronization**

Critical-Section Problem

Producer-Consumer Problem:

Producer Process: produces information

Consumer Process: consumes the information

- *unbounded-buffer* places no practical limit on the size of the buffer
- *bounded-buffer* assumes that there is a fixed buffer size

Critical-Section Problem

Shared data:

```
#define BUFFER_SIZE 10
typedef struct
{
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
int count = 0;
```

Critical-Section Problem

Bounded-Buffer:

Producer:

```
while (true)
{
    /* produce an item and
    put in nextProduced */

    while (count == BUFFER_SIZE)
        ; // do nothing
    buffer [in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
    count++;
}
```

Consumer:

```
while (true)
{
    while (count == 0)
        ; // do nothing
    nextConsumed = buffer[out];
    out = (out + 1) %
    BUFFER_SIZE;
    count--;

    /* consume the item in
    nextConsumed */
}
```

Count is used by both producer and consumer process

Critical-Section Problem

Race Condition:

- `count++` could be implemented as

```
register1 = count
register1 = register1 + 1
count = register1
```

- `count--` could be implemented as

```
register2 = count
register2 = register2 - 1
count = register2
```

- Consider this execution interleaving with “count = 5” initially:

```
S0: producer execute register1 = count {register1 = 5}
S1: producer execute register1 = register1 + 1 {register1 = 6}
S2: consumer execute register2 = count {register2 = 5}
S3: consumer execute register2 = register2 - 1 {register2 = 4}
S4: producer execute count = register1 {count = 6}
S5: consumer execute count = register2 {count = 4}
```

Each process is designed with an entry section followed by the critical section of that process. The exit section which follows the critical section is meant for relinquishing the right to modify the shared data. In the entry section, the process acquires the permission to modify the shared data. It is obvious that when a process executes in the CS, no other process will execute in their critical section.

Critical-Section Problem

```
do {
```

entry section

critical section

exit section

remainder section

```
} while (TRUE);
```

Every process will have its Own critical section involving the shared data/variable. It is not necessary that the critical section is similar for all processes, its only required that each critical section involves the shared data/variable.

Solution to the Critical-Section Problem

Solution must satisfy the following three conditions:

1. **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections.
 1. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely.
 1. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted.
- Assume that each process executes at a nonzero speed
 - No assumption concerning relative speed of the N processes

Peterson's Solution

- Two process solution
- Assume that the **LOAD** and **STORE** instructions are atomic; that is, cannot be interrupted.
- The two processes share two variables:
 - int **turn**;
 - Boolean **flag[2]**
- The variable **turn** indicates **whose turn** it is to enter the critical section.
- The **flag** array is used to indicate if a process is ready to enter the critical section.

flag[i] = true implies that process **P_i** is ready!

Peterson's Solution

Algorithm for Process P_i

```
while (true)
{
    flag[i] = TRUE;
    turn = j;

    while ( flag[j] && turn == j);

        CRITICAL SECTION

    flag[i] = FALSE;

        REMAINDER SECTION

}
```

Algorithm for Process P_j

```
while (true)
{
    flag[j] = TRUE;
    turn = i;

    while ( flag[i] && turn == i);

        CRITICAL SECTION

    flag[j] = FALSE;

        REMAINDER SECTION

}
```

Peterson's Solution

Algorithm for Process P_1

```
while (true)
{
    flag[1] = TRUE;
    turn = 2;

    while ( flag[2] && turn == 2);

    CRITICAL SECTION

    flag[1] = FALSE;

    REMAINDER SECTION

}
```

Algorithm for Process P_2

```
while (true)
{
    flag[2] = TRUE;
    turn = 1;

    while ( flag[1] && turn == 1);

    CRITICAL SECTION

    flag[2] = FALSE;

    REMAINDER SECTION

}
```

1. Mutual Exclusion 2. Progress 3. Bounded Waiting

Peterson's Solution

Mutual Exclusion is preserved:

Let us assume P_i execute in its critical section.

P_i can only execute in its CS only if

either $\text{flag}[j] = \text{false}$ or $\text{turn} = i$

If P_1 and P_2 could not have successfully executed their while statement because the value of turn can either be 1 or 2 but cannot be both.

Hence, only one process can execute the while statement successfully.

Say P_1 executes it successfully.

P_2 will spin on the while loop as long as the P_1 executes its CS.

Thus, mutual exclusion is preserved.

Multiple-Process Solutions

Bakery Algorithm:

Critical section for n processes:

- Before entering its critical section, process receives a number. Holder of the smallest number enters the critical section.
- If processes P_i and P_j receive the **same number**,
if $i < j$
 P_i is served first;
else
 P_j is served first.
- The numbering scheme always generates numbers in increasing order of enumeration; i.e., 1,2,3,3,3,3,4,5...

Multiple-Process Solutions

Bakery Algorithm:

- Notation: lexicographical order (ticket #, process id #)
 - $(a,b) < (c,d)$ if $a < c$ or if $a = c$ and $b < d$
 - $\max(a_0, \dots, a_{(n-1)})$ is a number, k , such that $k \geq a_i$ for $i = 0, \dots, n - 1$
- Shared data
 - var **choosing**: array $[0..n-1]$ of **boolean**;
 - number**: array $[0..n-1]$ of **integer**;
- Data structures are initialized to false and

Multiple-Process Solutions

Bakery Algorithm:

repeat

```
    choosing[i] := true;
    number[i] := max(number[0], number[1], ..., number[n - 1]) + 1;
    choosing[i] := false;
    for j := 0 to n - 1
        do begin
            while choosing[j] do no-op;
            while number[j] != 0
                and (number[j], j) < (number[i], i) do no-op;
        end;
```

critical section

```
number[i] := 0;
```

remainder section

until false;

Synchronization Hardware

- Many systems provide **hardware support** for critical section code
- **Uniprocessors** – could **disable interrupts**
 - Currently running code would execute without preemption
 - Generally too **inefficient on multiprocessor systems**
 - Operating systems using this **not** broadly **scalable**
- Modern machines provide special atomic hardware instructions
 - **Atomic = non-interruptable**
 - Either test memory word and set value
 - Or swap contents of two memory words

TestAndndSet Instruction

- **Definition:**

```
boolean TestAndSet (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

If two testAndSet instructions are simultaneously each on a different CPU, they are sequentially in some arbitrary order.

Solution using TestAndndSet

- Shared boolean variable lock., initialized to false.
- Solution:

```
while (true)
{
    while ( TestAndSet (&lock ))
        ; /* do nothing

        //  critical section

    lock = FALSE;

        //  remainder section

}
```

Solution using TestAndndSet

Definition:

```
boolean TestAndSet (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

Solution:

```
while (true)
{
    while ( TestAndSet (&lock ))
        ; /* do nothing

        // critical section

    lock = FALSE;

    // remainder section

}
```

Swap Instruction

- **Definition:**

```
void Swap (boolean *a, boolean *b)
{
    boolean temp = *a;
    *a = *b;
    *b = temp;
}
```

Solution using Swap

- Shared Boolean variable lock initialized to FALSE;
- Each process has a local Boolean variable key.
- Solution:

```
while (true)
```

```
{
```

```
    key = TRUE;  
    while ( key == TRUE)  
        Swap (&lock, &key );
```

```
        // critical section
```

```
        lock = FALSE;
```

```
        // remainder section
```

```
}
```

Solution using Swap

- **Definition:**

```
void Swap (boolean *a, boolean
*b)
{
    boolean temp = *a;
    *a = *b;
    *b = temp;
}
```

- Shared Boolean variable lock initialized to FALSE;
- Each process has a local Boolean variable key.

- **Solution:**

```
while (true)
```

```
{
```

```
    key = TRUE;
```

```
    while ( key == TRUE)
```

```
        Swap (&lock, &key );
```

```
        // critical section
```

```
        lock = FALSE;
```

```
        // remainder section
```

```
    }
```

Semaphore

- Synchronization tool that does not require **busy waiting**
- Semaphore S – **integer** variable
- Two standard operations modify S : **wait()** and **signal()**
 - Originally called **P()** and **V()**
- Less complicated
- Can only be accessed via two indivisible (atomic) operations
 - **wait (S)**

```
{  
    while S <= 0; // no-op  
    S--;  
}
```
 - **signal (S)**

```
{  
    S++;  
}
```


Semaphore as General Synchronization Tool

- **Counting semaphore** – integer value can range over an unrestricted domain
- **Binary semaphore** – integer value can range only between 0 and 1; can be **simpler to implement**
 - Also known as **mutex locks**
- **Can implement a counting semaphore S as a binary semaphore**
- Provides mutual exclusion
 - Semaphore S; // initialized to 1
 - wait (S);
 Critical Section
 signal (S);

Semaphore Implementation

- Must guarantee that no two processes can execute `wait ()` and `signal ()` on the same semaphore at the same time
- Thus, implementation becomes the critical section problem where the wait and signal code are placed in the critical section.
 - Could now have busy waiting in critical section implementation
 - But implementation code is short
 - Little busy waiting if critical section rarely occupied
- Note that applications may spend lots of time in critical sections and therefore this is not a good solution.

Spin lock

- Software solution and the semaphore definition given here **involves busy waiting**.
- This type of semaphore which involves busy waiting is called **spinlock**.
- A spinlock avoids context switch hence useful **in case of short critical section**.
- Useful in multiprocessors, **a thread busy-waits on a resource on one processor while another process uses the resource on a different processor**.

Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated **waiting queue**. Each entry in a waiting queue has two data items:
 - value (of type integer)
 - pointer to next record in the list
- Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue.
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue.

Semaphore Implementation with no Busy waiting (Cont.)

- Implementation of wait:

```
wait (S)
{
    value--;
    if (value < 0)
    {
        add this process to waiting queue
        block();
    }
}
```

- Implementation of signal:








```
Signal (S)
{
    value++;
    if (value <= 0) {
        remove a process P from waiting queue
        wakeup(P); }
}
```

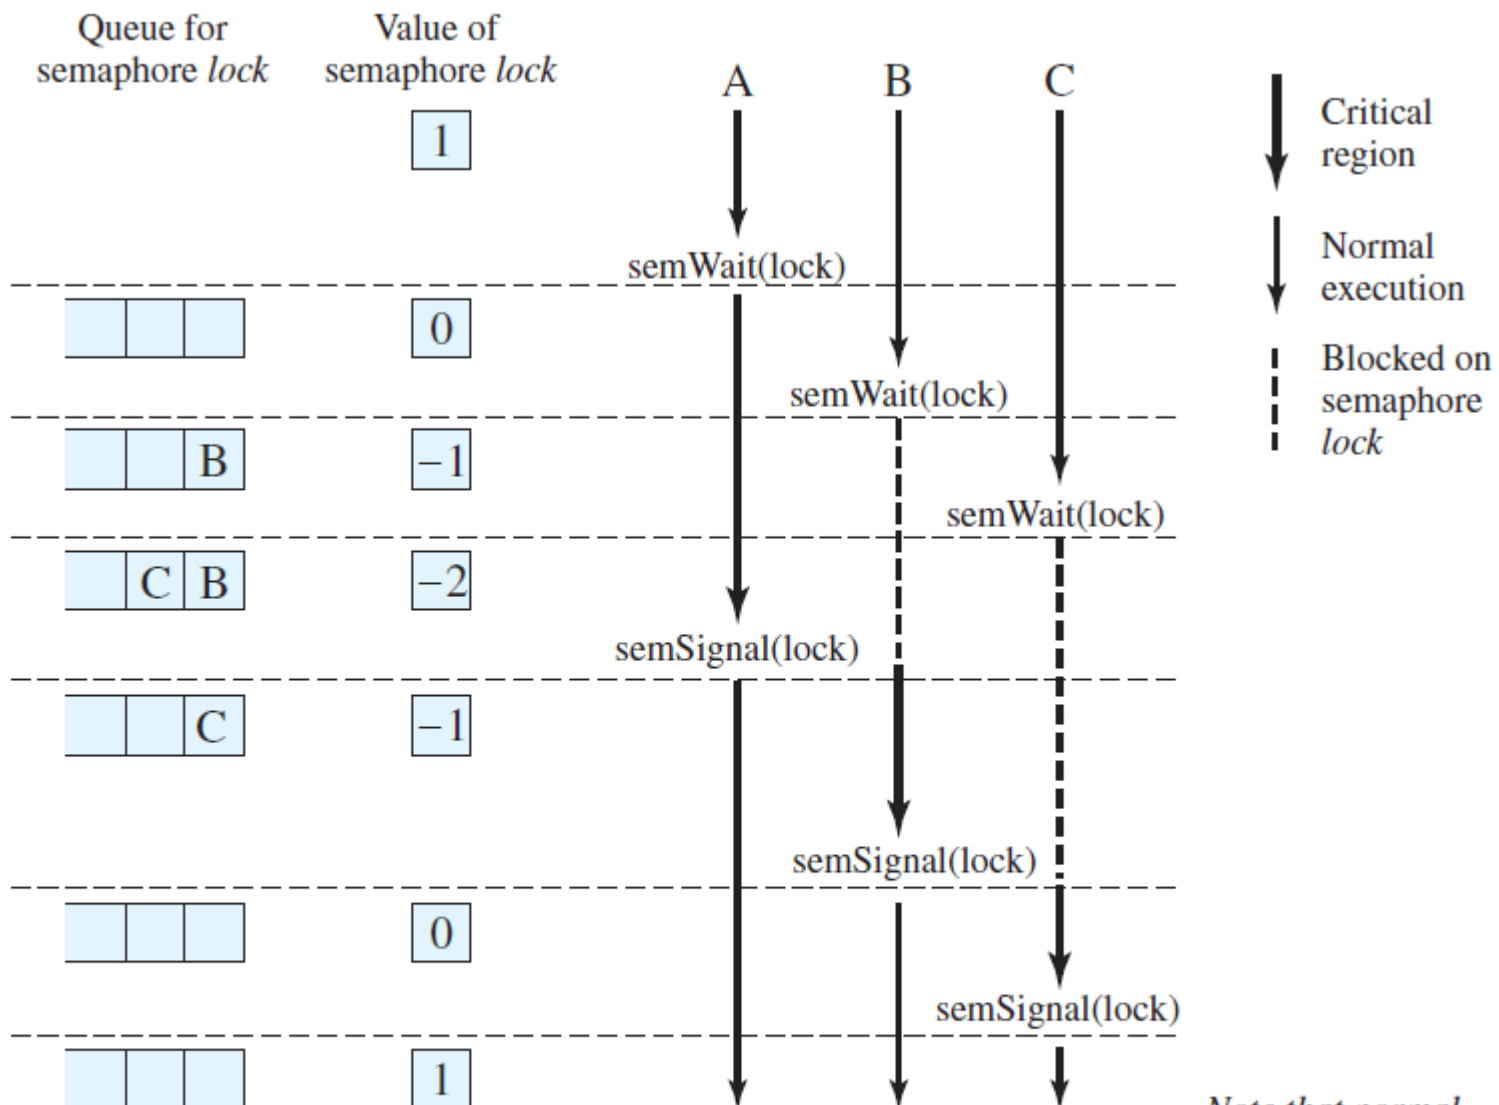
Binary Semaphore Definitions

```
struct binary_semaphore {
    enum {zero, one} value;
    queueType queue;
};

void semWaitB(binary_semaphore s)
{
    if (s.value == one)
        s.value = zero;
    else {
        /* place this process in s.queue */;
        /* block this process */;
    }
}

void semSignalB(semaphore s)
{
    if (s.queue is empty())
        s.value = one;
    else {
        /* remove a process P from s.queue */;
        /* place process P on ready list */;
    }
}
```





Note that normal execution can proceed in parallel but that critical regions are serialized.

Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes.

- Let **S** and **Q** be two semaphores initialized to 1

P_0
wait (S);
wait (Q);
.
.
.
signal (S);
signal (Q);

P_1
wait (Q);
wait (S);
.
.
.
signal (Q);
signal (S);

- **Starvation** – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.

Classical Problems of Synchronization

- **Bounded-Buffer Problem**
- **Readers and Writers Problem**
- **Dining-Philosophers Problem**

Bounded-Buffer Problem

- N buffers, each can hold one item
- Semaphore **mutex** initialized to the value 1
- Semaphore **full** initialized to the value 0
- Semaphore **empty** initialized to the value N .

Bounded Buffer Problem (Cont.)

- The structure of the producer process

```
while (true) {  
  
    // produce an item  
  
    wait (empty);  
    wait (mutex);  
  
    // add the item to the buffer  
  
    signal (mutex);  
    signal (full);  
}
```

Bounded Buffer Problem (Cont.)

- The structure of the consumer process

```
while (true)
{
    wait (full);
    wait (mutex);

    // remove an item from buffer

    signal (mutex);
    signal (empty);

    // consume the removed item
}
```

Bounded Buffer Problem (Cont.)

Producer:

```
while (true)
{
    // produce an item

    wait (empty);
    wait (mutex);

    // add the item to the
buffer

    signal (mutex);
    signal (full);
}
```

Consumer:

```
while (true)
{
    wait (full);
    wait (mutex);
    //remove an item from
buffer

    signal (mutex);
    signal (empty);
    // consume the
removed item
}
```

```
int n;  
binary_semaphore s = 1, delay = 0;  
void producer()  
{  
    while (true) {  
        produce();  
        semWaitB(s);  
        append();  
        n++;  
        if (n==1) semSignalB(delay);  
        semSignalB(s);  
    }  
}
```

```
void consumer()  
{  
    int m; /* a local variable */  
    semWaitB(delay);  
    while (true) {  
        semWaitB(s);  
        take();  
        n--;  
        m = n;  
        semSignalB(s);  
        consume();  
        if (m==0) semWaitB(delay);  
    }  
}
```

Solution to the Infinite-Buffer Producer/Consumer Problem

```
semaphore n = 0, s = 1;
void producer()
{
    while (true) {
        produce();
        {
            semWait(s);
            append();
            semSignal(s);
        }
        semSignal(n);
    }
}
```

```
void consumer()
{
    while (true) {
        semWait(n);
        {
            semWait(s);
            take();
            semSignal(s);
        }
        consume();
    }
}
```

Solution to the Infinite-Buffer Producer/Consumer Problem

Readers-Writers Problem

- A data set is shared among a number of concurrent processes
 - **Readers** – only read the data set; they do **not** perform any updates
 - **Writers** – can both read and write.
- Problem – allow multiple readers to read at the same time. Only one single writer can access the shared data at the same time.
- Shared Data
 - **Data set**
 - Semaphore **mutex** initialized to 1.
 - Semaphore **wrt** initialized to 1.
 - Integer **readcount** initialized to 0.

Readers-Writers Problem (Cont.)

- The structure of a writer process

```
while (true)
{
    wait (wrt) ;

    //  writing is performed

    signal (wrt) ;
}
```

Readers-Writers Problem (Cont.)

- The structure of a reader process

```
while (true) {  
    wait (mutex) ;  
    readcount ++ ;  
    if (readcount == 1) wait (wrt) ;  
    signal (mutex);  
  
    // reading is performed  
  
    wait (mutex) ;  
    readcount - - ;  
    if (readcount == 0) signal (wrt) ;  
    signal (mutex) ;  
}
```

Readers-Writers Problem (Cont.)

Writer process

```
while (true)
{
    wait (wrt) ;

    // writing is performed

    signal (wrt) ;
}
```

Reader process

```
while (true)
{
    wait (mutex) ;
    readcount ++ ;
    if (readcount == 1)
        wait (wrt) ;
    signal (mutex)

    // reading is performed

    wait (mutex) ;
    readcount -- ;
    if (readcount == 0)
        signal (wrt) ;
    signal (mutex) ;
}
```

Writers have priority:

No new readers are allowed once at least one writer has declared the desire to write.

Writers have priority:

- A semaphore **rsem** that inhibits all readers while there is at least one writer desiring access to the data area
- A variable **writecount** that controls the setting of rsem
- A semaphore **y** that controls the updating of writecount

Readers/Writers w/ Writers Priority (Using Semaphores)

Reader:

```
P(mutex);
if (AW+WW > 0)
    WR++;
else {
    V(OKToRead);
    AR++;
}
V(mutex);
P(OKToRead);

read database

P(mutex);
AR--;
if (AR == 0 &&
    WW > 0) {
    V(OKToWrite);
    AW++; WW--;
}
V(mutex);
```

Writer:

```
P(mutex);
if (AW+AR > 0)
    WW++;
else {
    V(OKToWrite);
    AW++;
}
V(mutex);
P(OKToWrite);

write database

P(mutex);
AW--;
if (WW > 0) {
    V(OKToWrite);
    AW++; WW--;
} else if (WR > 0) {
    V(OKToRead);
    AR++; WR--;
}
V(mutex);
```

Notes on R/W w/ Writers Priority (Using Semaphores)


■ Reader:

- If there are active or waiting writers, this reader has to wait (writers have priority)
- Otherwise, this reader can read (possibly along with other readers)
- When the last reader finishes, if there are waiting writers, it must wake one up

■ Writer:

- If there are active readers or writers, this writer has to wait (everyone has to finish before writer can update database)
- Otherwise, this writer can write (and has exclusive access to database)
- When the writer finishes,
 - (first choice) if there are waiting writers, it must wake one up (writers have priority)
 - (second choice) if there are waiting readers, it must wake one up

Dining-Philosophers Problem



- Shared data
 - Bowl of rice (data set)
 - Semaphore **chopstick [5]** initialized to 1

Dining-Philosophers Problem (Cont.)

- The structure of Philosopher *i*:

```
while (true)
```

```
{
```

```
    wait ( chopstick[i] );
```

```
    wait ( chopStick[ (i + 1) % 5] );
```

```
        // eat
```

```
    signal ( chopstick[i] );
```

```
    signal ( chopstick[ (i + 1) % 5] );
```

```
        // think
```

```
}
```

End of Session 5

Shared Variable in Threaded Programs

Multiple threads can share the same program variables.

- 1 #include "csapp.h"
- 2 #define N 2
- 3
- 4 char **ptr; /* global variable */
- 5
- 6 void *thread(void *vargp);
- 7
- 8 int main()
- 9 {
- 10 int i;
- 11 pthread_t tid;
- 12 char *msgs[N] = {
- 13 "Hello from foo",
- 14 "Hello from bar"
- 15 };
- 16
- 17 ptr = msgs;
- 18
- 19 for (i = 0; i < N; i++)
- 20 pthread_create(&tid, NULL, thread, (void *)i);
- 21 pthread_exit(NULL);
- 22 }
- 23
- 24 void *thread(void *vargp)
- 25 {
- 26 int myid = (int)vargp;
- 27 static int cnt = 0;
- 28
- 29 printf("[%d]: %s (cnt=%d)\n", myid, ptr[myid], ++cnt);
- 30 }

- **unix> ./sharing**
- **[0]: Hello from foo (cnt=1)**
- **[1]: Hello from bar (cnt=2)**

Threads Memory Model

A pool of concurrent threads runs in the context of a process.

Each thread has its own separate thread context, which includes a **thread ID**, **stack**, **stack pointer**, **program counter**, **condition codes**, and **general purpose register values**.

Threads Memory Model

Each thread shares the rest of the process context with the other threads. This includes the entire user virtual address space, which consists of read-only text (code), read/write data, the heap, and any shared library code and data areas.

Threads Memory Model

The threads also share the **same set of open files** and **the same set of installed signal handlers**.

In an operational sense, it is impossible for one thread to read or write the register values of another thread.

C variables in threaded programs are mapped to virtual memory according to their storage classes.

Global variables.

A *global variable* is any variable **declared outside** of a function. At run-time, the **read/write area of virtual memory contains exactly one instance of each global variable that can be referenced by any thread.**

Global variables.

For example, the global ptr variable in line 4 has one run-time instance in the read/write area of virtual memory.

When there is only one instance of a variable, we will denote the instance by simply using the variable name, in this case ptr.

Local automatic variables.

A local automatic variable is one that is declared **inside a function** without the static attribute. At run-time, each thread's stack contains its own instances of any local automatic variables. **This is true even if multiple threads execute the same thread routine.**

Local automatic variables.

For example, there is one instance of the local variable `tid`, and it resides on the stack of the main thread. We will denote this instance as `tid.m`. As another example, there are two instances of the local variable `myid`, one instance on the stack of peer thread 0, and the other on the stack of peer thread 1. We will denote these instances as `myid.p0` and `myid.p1` respectively.

Mapping Variables to Memory

Local static variables.

A *local static variable* is one that is declared **inside a function with the static attribute**. As with global variables, the **read/write area of virtual memory contains exactly one instance of each local static variable declared in a program**.

For example, even though each peer thread in our example program declares **cnt** in line 27, at runtime there is only one instance of **cnt** residing in the **read/write area of virtual memory**. Each peer thread reads and writes this instance.

Mapping Variables to Memory

Shared Variables

A variable `v` is shared if and only if one of its instances is referenced by more than one thread.

For example, variable `cnt` in our example program is shared because it has only one run-time instance, and this instance is referenced by both peer threads.

On the other hand, `myid` is not shared because each of its two instances is referenced by exactly one thread.

However, it is important to realize that local automatic variables such as `msgs` can also be shared.

Progress Graphs

A *progress graph* models the execution of n concurrent threads as a trajectory through an n -dimensional Cartesian space.

Each **axis k** corresponds to the progress of **thread k** .

Each point $(I_1; I_2; \dots; I_n)$ represents the state where thread k , ($k = 1; \dots; n$) has **completed instruction I_k** .

The **origin** of the graph corresponds to the *initial state* where none of the threads has yet completed an instruction.


```
1 #include "csapp.h"
2
3 #define NITERS 1000000000
4
5 void *count(void *arg);
6
7 /* shared variable */
8 unsigned int cnt = 0;
9
10 int main()
11 {
12     pthread_t tid1, tid2;
13
14     pthread_create(&tid1, NULL, count, NULL);
15     pthread_create(&tid2, NULL, count, NULL);
16
17     pthread_join(tid1, NULL);
18     pthread_join(tid2, NULL);
19
20     if (cnt != (unsigned)NITERS*2)
21         printf("BOOM! cnt=%d\n", cnt);
22     else
23         printf("OK cnt=%d\n", cnt);
24     exit(0);
25 }
26
27 /* thread routine */
28 void *count(void *arg)
29 {
30     int i;
31
32     for (i=0; i<NITERS; i++)
33         cnt++;
34     return NULL;
35 }
```

C code for thread i

```
for (i=0; i<NITERS; i++)  
    ctr++;
```



Asm code for thread i

<pre>.L9: movl -4(%ebp), %eax cmpl \$99999999, %eax jle .L12 jmp .L10</pre>	} H_i : Head
<pre>.L12: movl ctr, %eax leal 1(%eax), %edx movl %edx, ctr</pre>	
<pre>.L11: movl -4(%ebp), %eax leal 1(%eax), %edx movl %edx, -4(%ebp) jmp .L9 .L10:</pre>	} T_i : Tail

L_i : Load ctr
 U_i : Update ctr
 S_i : Store ctr

- H_i : The block of instructions at the head of the loop.
- L_i : The instruction that loads the shared variable `cnt` into register $\%eax_i$, where $\%eax_i$ denotes the value of register $\%eax$ in thread i .
- U_i : The instruction that updates (increments) $\%eax_i$.
- S_i : The instruction that stores the updated value of $\%eax_i$ back to the shared variable `cnt`.
- T_i : The block of instructions at the tail of the loop.

```
unix> ./badcnt  
BOOM! ctr=198841183
```

```
unix> ./badcnt  
BOOM! ctr=198261801
```

```
unix> ./badcnt  
BOOM! ctr=198269672
```

Instructions can be interleaved in any order, so long as the instructions for each thread execute in program order.

For example, the ordering

H1;H2; L1; L2; U1; U2; S1; S2; T1; T2

is sequentially consistent, while the ordering

H1;H2; U1; L2; L1; U2; S1; S2; T1; T2

is not sequentially consistent because U1 executes before L1.

Unfortunately not all sequentially consistent orderings are created equal. Some will produce correct results, but others will not, and there is no way for us to predict whether the operating system will choose a correct ordering for our threads.




Figure 11.11: Progress graph for the first loop iteration of `badcnt.c`.

The horizontal axis corresponds to thread 1,
the vertical axis to thread 2.

Point $(L_1; S_2)$ corresponds to the state where thread 1 has completed L_1 and thread 2 has completed S_2 .




Figure 11.12: Legal transitions in a progress graph.

For the **single-processor** systems that we are concerned about, where instructions complete one at a time in **sequentially-consistent** order, legal transitions move to the **right** (an instruction in thread 1 completes) or **up** (an instruction in thread 2 completes).

Programs never run backwards, so transitions that move down or to the left are not legal.




Figure 11.13: An example trajectory.

$H_1, L_1, U_1, H_2, L_2, S_1, T_1, U_2, S_2, T_2$.

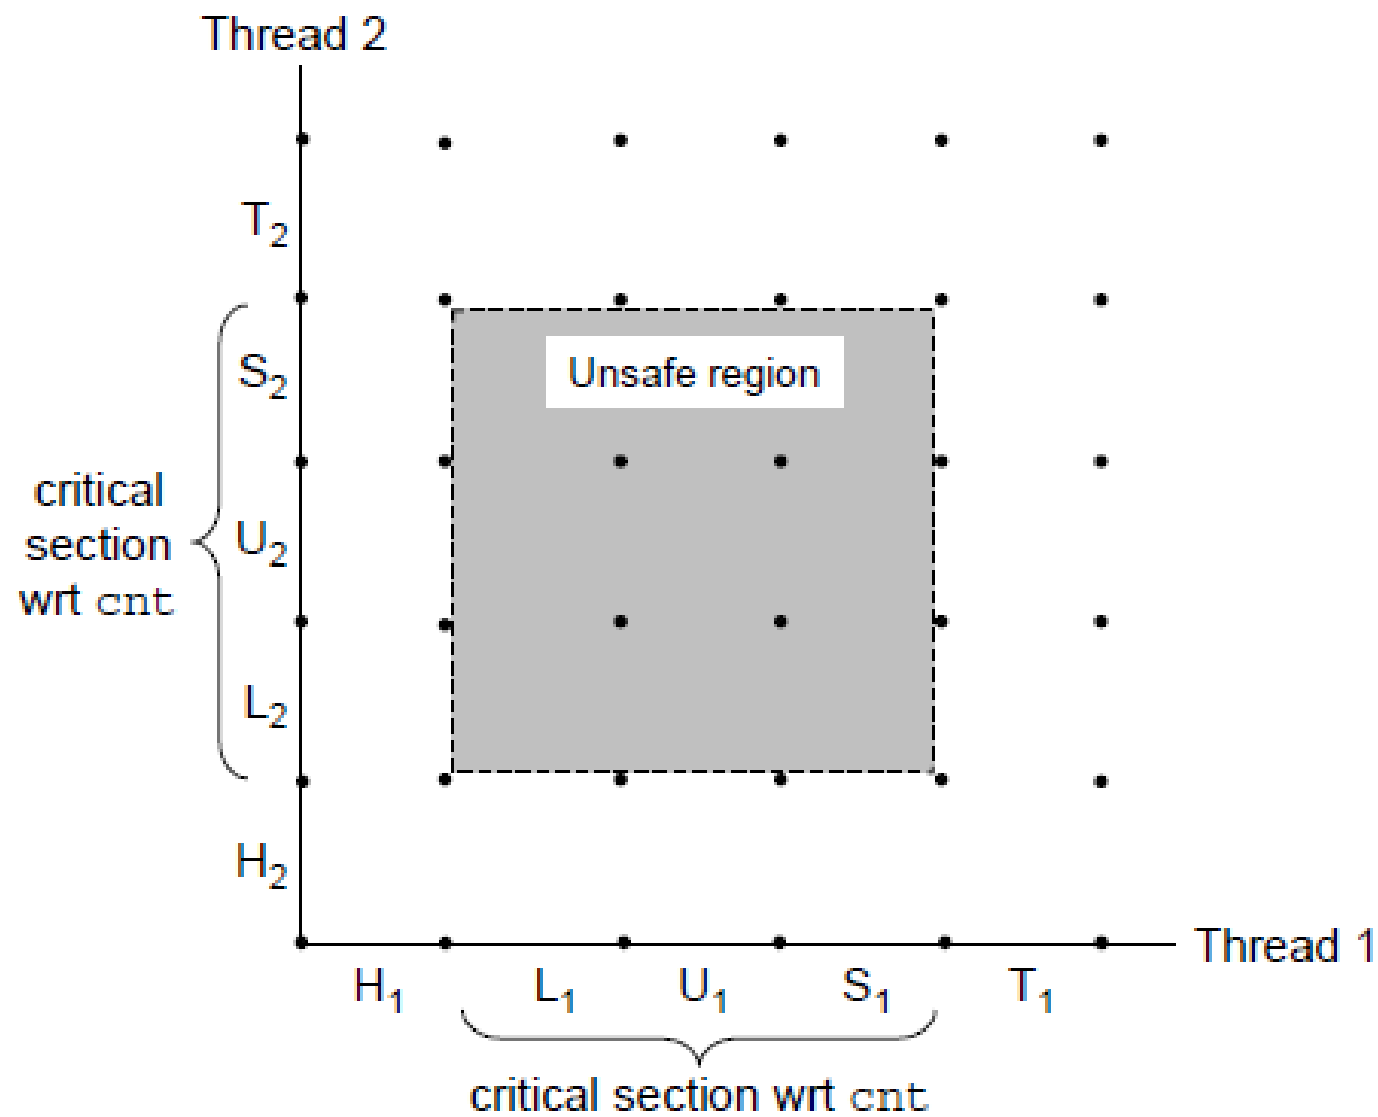


Figure 11.14: **Critical sections and unsafe regions.**

For thread i , the instructions $(L_i; U_i; S_i)$ that manipulate the contents of the shared variable cnt constitute a *critical section* (with respect to shared variable cnt) that should not be interleaved with the critical section of the other thread.

The intersection of the two critical sections defines a region of the state space known as an *unsafe region*.

Figure 11.14 shows the unsafe region for the variable cnt . Notice that the unsafe region abuts, but does not include, the states along its perimeter. For example, states $(H1; H2)$ and $(S1; U2)$ abut the unsafe region, but are not a part of it.


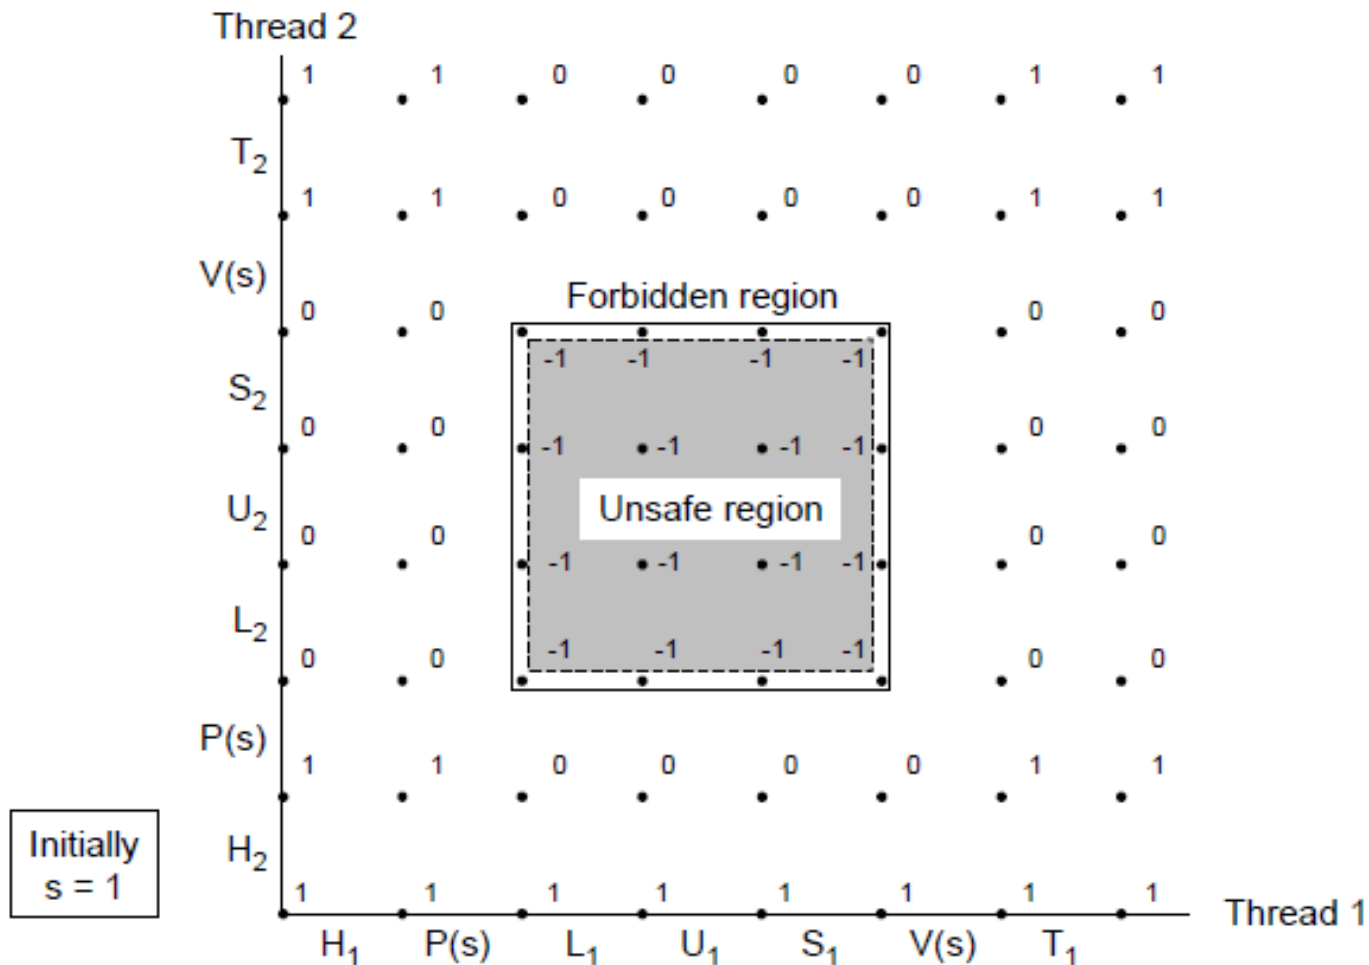



Figure 11.15: Safe and unsafe trajectories.

The basic idea is to associate a semaphore s , initially 1, with each shared variable (or related set of shared variables) and then surround the corresponding critical section with $P(s)$ and $V(s)$ operations.





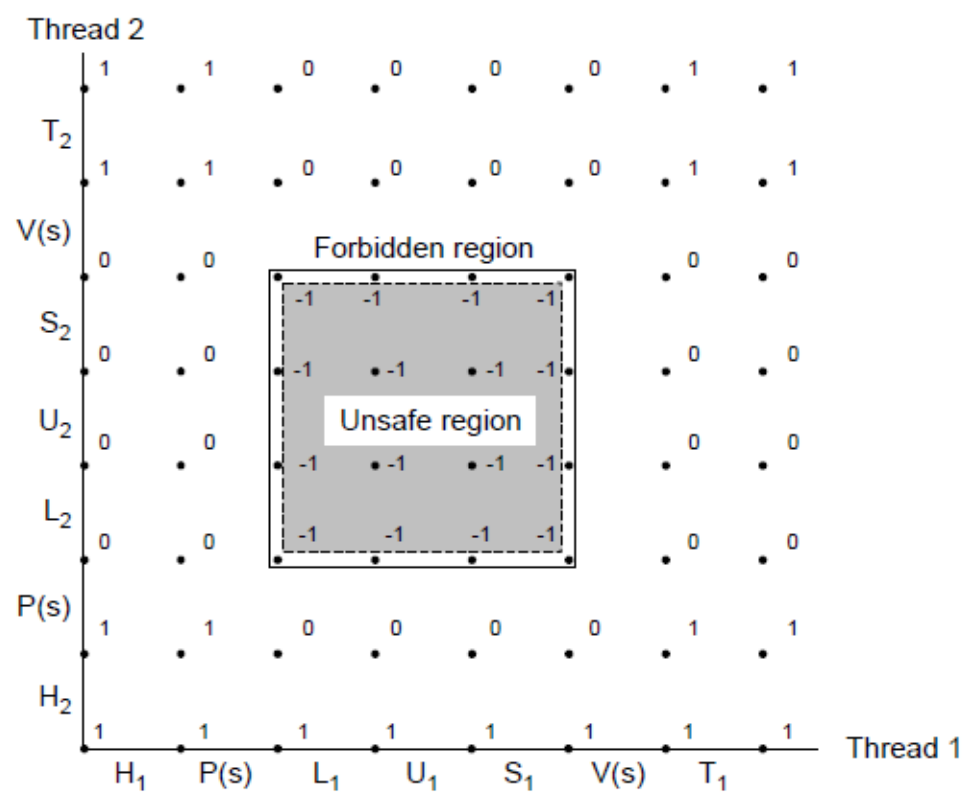



Figure 11.15: Safe and unsafe trajectories.

In the figure, each state is labeled with the value of semaphore s in that state.

The crucial idea is that this combination of P and V operations creates a collection of states, called a *forbidden region*, where $s < 0$.

Because of the semaphore invariant, no feasible trajectory can include one of the states in the forbidden region.

And since the forbidden region completely encloses the unsafe region, no feasible trajectory can touch any part of the unsafe region.

Thus, *every feasible trajectory is safe*, and regardless of the ordering of the instructions at runtime, the program correctly increments the counter.

A primary aim of an operating system is **to share** a computer installation among many programs making unpredictable demands upon its resources.

Designer should try to construct **separate schedulers** for each class of resource.

A primary task of its designer is therefore to construct resource allocation (or scheduling) algorithms for resources of various kinds (main store, drum store, magnetic tape handlers, consoles, etc.).

Each scheduler will consist of a certain amount of local administrative data, together with some procedures and functions which are called by programs wishing to acquire and release resources.

Such a collection of associated data and procedures is known as a *monitor*.

Monitor

A monitor is an object intended to be used safely by more than one thread.

Monitor

Its methods are executed with mutual exclusion.

At each point in time, at most one thread may be executing any of its methods.

Monitor

Monitors also provide a mechanism for threads to temporarily **give up exclusive access**, in order to wait for some **condition to be met**, before regaining exclusive access and resuming their task.

Monitor

Monitors also have a mechanism for signaling other threads that such conditions have been met.

Mutual Exclusion

While a thread is executing a method of a monitor, it is said to *occupy* the monitor.

Mutual exclusion property:
at each point in time, at most one thread may occupy the monitor.

Upon calling one of the methods, a thread must wait until no thread is executing any of the monitor's methods before starting execution of its method.

In a simple implementation, mutual exclusion can be implemented by the compiler equipping each monitor object with a private lock, often in the form of a semaphore.

This lock is initially unlocked, is locked at the start of each public method, and is unlocked at each return from each public method.

monitor class *Account*

{

private *int* balance := 0

invariant balance >= 0

public method *boolean* withdraw(*int* amount)

{

if amount < 0 then error "Amount may not be negative"

else if balance < amount then return false

else { balance := balance - amount ; return true }

}

public method deposit(*int* amount)

{

if amount < 0 then error "Amount may not be negative"

else balance := balance + amount

}

}

Waiting and Signaling

For many applications, mutual exclusion is not enough.

Threads attempting an operation may need to wait until some assertion P holds true.

A busy waiting loop

`while not (P) do skip`

will not work, as mutual exclusion will prevent any other thread from entering the monitor to make the condition true.

Waiting and Signaling

Condition variables:

A condition variable is a queue of threads, associated with a monitor, upon which a thread may wait for some assertion to become true.

Thus each condition variable c is associated with some assertion P_c .

While a thread is waiting upon a condition variable, that thread is not considered to occupy the monitor, and so other threads may enter the monitor to change the monitor's state.

In most types of monitors, these other threads may signal the condition variable c to indicate that assertion P_c is true.

Two main operations on conditions variables:

wait c is called by a thread that needs to wait until the assertion P_c to be true before proceeding.

signal c (sometimes written as **notify c**) is called by a thread to indicate that the assertion P_c is true.

```

monitor class Semaphore
{
    private int s := 0
    invariant s >= 0
    private Condition sIsPositive /* associated with s > 0 */

    public method P()
    {
        if s = 0 then wait sIsPositive assert s > 0 s := s - 1
    }

    public method V()
    {
        s := s + 1 assert s > 0
        signal sIsPositive
    }
}

```

A thread that tries to decrement must wait until the integer is positive. We use a condition variable *sIsPositive* with an associated assertion of *PsIsPositive* = $(s > 0)$


When a signal happens on a condition that at least one other thread is waiting on, there are at least two threads that could then occupy the monitor:

- the thread that signals and
- any one of the threads that is waiting.

In order that at most one thread occupies the monitor at each time, a choice must be made. Two schools of thought exist on how best to resolve this choice. This leads to two kinds of condition variables which will be examined next:

Blocking condition variables give priority to a signaled thread.

Nonblocking condition variables give priority to the signaling thread.



enter the monitor:

enter the method

if the monitor is locked

add this thread to e block this thread

else

lock the monitor

leave the monitor:

schedule

return from the method

wait c :

add this thread to c.q

schedule

block this thread

schedule :

if there is a thread on **s**

select and remove one thread from s and restart it

(this thread will occupy the monitor next)

else if there is a thread on e

select and remove one thread from e and restart it

(this thread will occupy the monitor next)

else

unlock the monitor

(the monitor will become unoccupied)

signal c :

if there is a thread waiting on c.q

select and remove one such

thread t from c.q

(t is called "the signaled thread")

add this thread to s restart t

(so t will occupy the monitor

next)

block this thread

With a blocking condition variable, the signaling thread must wait outside the monitor (at least) until the signaled thread relinquishes occupancy of the monitor by either returning or by again waiting on a condition.


We assume there are **two queues** of threads associated with each monitor object

- e is the entrance queue

- s is a queue of threads that have signaled.

In addition we assume that for each condition c , there is a **queue**

$c.q$, which is a queue for threads waiting on condition c



With *nonblocking condition variables*, signaling does not cause the signaling thread to lose occupancy of the monitor. Instead the signaled threads are moved to the `e` queue. Signaling thread continues. There is no need for the `s` queue.

enter the monitor:

- enter the method
- if the monitor is locked
 - add this thread to e
 - block this thread
- else
 - lock the monitor

leave the monitor:

- schedule
- return from the method

notify all c :

- move all threads waiting on c.q to e

schedule :

- if there is a thread on e
 - select and remove one
 - thread from e and restart it
- else
 - unlock the monitor

wait c :

- add this thread to c.q
- schedule
- block this thread

notify c :

- if there is a thread waiting on c.q
- select and remove one thread t from c.q
- (t is called "the notified thread")
- move t to e

Monitors

A monitor is an object that contains both
the **data** and
the **procedures**
needed to perform allocation of a particular of a serially reusable
shared resources.

A thread calls a **monitor entry routine** to access a resource.

Only **one** thread **at a time** is allowed to enter the monitor.

Other threads are made to wait at the **monitor boundary**.

Data inside a monitor may be
i) **global to all routines inside the monitor, or**
ii) **local to a specific routine**

Monitors

Data inside a monitor may be

i) global to all routines inside the monitor, or

ii) local to a specific routine

Monitor data is accessible only within the monitor.

Threads outside the monitor cannot access the monitor data.

This is a kind of **information hiding**.

Monitors

A thread –

- a) calls a monitor entry routine**
- b) if no other threads inside the monitor**
- c) the thread acquires a lock on the monitor and enters it**
- d) else**
- e) the thread is made to wait until the lock is released by the**
other thread

Monitors

Finally, a thread having the lock of the monitor calls the monitor entry routine to release the resource.

Monitor entry routine calls signal to allow one of the waiting threads to enter the monitor and acquire the resource.

**If there is no waiting thread then signal has no effect,
monitor recaptures the resource.**

Monitors : Condition variables

Condition variables:

A thread inside a monitor uses a condition variable to wait on a condition outside the monitor.

A monitor associates a separate condition variable with each distinct situation that might cause a thread to have to wait.

Operations:

wait(condition variable)

signal(condition variable)

Monitors : Condition variables

Condition variables ... :

Every condition variable has an associated Queue.

A thread calls wait on a c.v. is placed in the queue of the c.v.

**→ While in queue the thread is waiting outside the monitor.
(so that another thread may enter the monitor to signal.)**

A thread calls signal on a c.v. causes a thread waiting in the queue to be removed and reenter the monitor.

→ FIFO queue is maintained most often.

Monitors : Condition variables

Condition variables ... :

Signal-and-exit monitor:

- a thread immediately exits the monitor upon signaling.

Signal-and-continue monitor:

- signals that the monitor will soon be available
- still keeps the lock until the thread exits
 - a) calling a wait on another c.v.
 - b) after executing some code in the monitor

Monitors : Resource Allocation

Resource Allocation:

```
1 // Fig. 6.1: Resource allocator monitor
2
3 // monitor initialization (performed only once)
4 boolean inUse = false; // simple state variable
5 Condition available; // condition variable
6
7 // request resource
8 monitorEntry void getResource()
9 {
10     if ( inUse ) // is resource in use?
11     {
12         wait( available ); // wait until available is signaled
13     } // end if
14
15     inUse = true; // indicate resource is now in use
16
17 } // end getResource
18
19 // return resource
20 monitorEntry void returnResource()
21 {
22     inUse = false; // indicate resource is not in use
23     signal( available ); // signal a waiting thread to proceed
24
25 } // end returnResource
```

Monitors : Circular buffer

```
1  // Fig. 6.2: Circular buffer monitor
2
3  char circularBuffer[] = new char[ BUFFER_SIZE ]; // buffer
4  int writerPosition = 0; // next slot to write to
5  int readerPosition = 0; // next slot to read from
6  int occupiedSlots = 0; // number of slots with data
7  Condition hasData;      // condition variable
8  Condition hasSpace;     // condition variable
9
10 // monitor entry called by producer to write data
11 monitorEntry void putChar( char slotData )
12 {
13     // wait on condition variable hasSpace if buffer is full
14     if ( occupiedSlots == BUFFER_SIZE )
15     {
16         wait( hasSpace ); // wait until hasSpace is signaled
17     } // end if
18
19     // write character to buffer
20     circularBuffer[ writerPosition ] = slotData;
21     ++occupiedSlots; // one more slot has data
22     writerPosition = (writerPosition + 1) % BUFFER_SIZE;
23     signal( hasData ); // signal that data is available
24 } // end putChar
25
```


Monitors : Circular buffer

```
26 // monitor entry called by consumer to read data
27 monitorEntry void getChar( outputParameter slotData )
28 {
29     // wait on condition variable hasData if the buffer is empty
30     if ( occupiedSlots == 0 )
31     {
32         wait( hasData ); // wait until hasData is signaled
33     } // end if
34
35     // read character from buffer into output parameter slotData
36     slotData = circularBuffer[ readPosition ];
37     occupiedSlots--; // one fewer slots has data
38     readerPosition = (readerPosition + 1) % BUFFER_SIZE;
39     signal( hasSpace ); // signal that character has been read
40 } // end getChar
```

Monitors : Readers Writers Problem

```
1 // Fig. 6.3: Readers/writers problem
2
3 int readers = 0; // number of readers
4 boolean writeLock = false; // true if a writer is writing
5 Condition canWrite; // condition variable
6 Condition canRead; // condition variable
7
8 // monitor entry called before performing read
9 monitorEntry void beginRead()
10 {
11     // wait outside monitor if writer is currently writing or if
12     // writers are currently waiting to write
13     if ( writeLock || queue( canWrite ) )
14     {
15         wait( canRead ); // wait until reading is allowed
16     } // end if
17
18     ++readers; // there is another reader
19
20     signal( canRead ); // allow waiting readers to proceed
21 } // end beginRead
22
```


Monitors : Readers Writers Problem

```
23 // monitor entry called after reading
24 monitorEntry void endRead()
25 {
26     --readers; // there are one fewer readers
27
28     // if no more readers are reading, allow a writer to write
29     if ( readers == 0 )
30     {
31         signal ( canWrite ); // allow a writer to proceed
32     } // end if
33
34 } // end endRead
35
36 // monitor entry called before performing write
37 monitorEntry void beginWrite()
38 {
39     // wait if readers are reading or if a writer is writing
40     if ( readers > 0 || writeLock )
41     {
42         wait( canWrite ); // wait until writing is allowed
43     } // end if
44 }
```

Monitors : Readers Writers Problem

```
45     writeLock = true; // lock out all readers and writers
46 } // end beginWrite
47
48 // monitor entry called after performing write
49 monitorEntry void endWrite()
50 {
51     writeLock = false; // release lock
52
53     // if a reader is waiting to enter, signal a reader
54     if ( queue( canRead ) )
55     {
56         signal( canRead ); // cascade in waiting readers
57     } // end if
58     else // signal a writer if no readers are waiting
59     {
60         signal( canWrite ); // one waiting writer can proceed
61     } // end else
62
63 } // end endWrite
```


Java Multithreading Case Study, Part III: Producer/Consumer Relationship in Java

```
1 // Fig. 6.4: SynchronizedBuffer.java
2 // SynchronizedBuffer synchronizes access to a shared integer.
3
4 public class SynchronizedBuffer implements Buffer
5 {
6     private int buffer = -1; // shared by producer and consumer
7     private int occupiedBuffers = 0; // counts occupied buffers
8
9     // place value into buffer
10    public synchronized void set( int value )
11    {
12        // for display, get name of thread that called this method
13        String name = Thread.currentThread().getName();
14
15        // while no empty buffers, place thread in waiting state
16        while ( occupiedBuffers == 1 )
17        {
18            // output thread and buffer information, then wait
19            try
20            {
21                System.err.println( name + " tries to write." );
22                displayState( "Buffer full. " + name + " waits." );
23                wait(); // wait until buffer is empty
24            } // end try
25        }
```

Java Multithreading Case Study, Part III: Producer/Consumer Relationship in Java

```
26         // if waiting thread interrupted, print stack trace
27         catch ( InterruptedException exception )
28         {
29             exception.printStackTrace();
30         } // end catch
31
32     } // end while
33
34     buffer = value; // set new buffer value
35
36     // indicate producer cannot store another value
37     // until consumer retrieves current buffer value
38     ++occupiedBuffers;
39
40     displayState( name + " writes " + buffer );
41
42     notify(); // tell waiting thread to enter ready state
43 } // end method set; releases lock on SynchronizedBuffer
44
```

Java Multithreading Case Study, Part III: Producer/Consumer Relationship in Java

```
44
45     // return value from buffer
46     public synchronized int get()
47     {
48         // for display, get name of thread that called this method
49         String name = Thread.currentThread().getName();
50
51         // while no data to read, place thread in waiting state
52         while ( occupiedBuffers == 0 )
53         {
54             // output thread and buffer information, then wait
55             try
56             {
57                 System.err.println( name + " tries to read." );
58                 displayState( "Buffer empty. " + name + " waits." );
59
60                 wait();// wait until buffer contains new values
61             } // end try
62
63             // if waiting thread interrupted, print stack trace
64             catch ( InterruptedException exception )
65             {
66                 exception.printStackTrace();
67             } // end catch
68         } // end while
69     }
```


Java Multithreading Case Study, Part III: Producer/Consumer Relationship in Java

```
70      // indicate that producer can store another value
71      // because consumer just retrieved buffer value
72      --occupiedBuffers;
73
74      displayState( name + " reads " + buffer );
75
76      notify(); // tell waiting thread to become ready
77
78      return buffer;
79  } // end method get; releases lock on SynchronizedBuffer
80
81  // display current operation and buffer state
82  public void displayState( String operation )
83  {
84      StringBuffer outputLine = new StringBuffer( operation );
85      outputLine.setLength( 40 );
86      outputLine.append( buffer + "\t\t" + occupiedBuffers );
87      System.err.println( outputLine );
88      System.err.println();
89  } // end method displayState
90
91  } // end class SynchronizedBuffer
```

6.4 Java Multithreading Case Study, Part III: Producer/Consumer Relationship in Java

```
1 // Fig. 6.5: SharedBufferTest2.java
2 // SharedBufferTest2 creates producer and consumer threads.
3
4 public class SharedBufferTest2
5 {
6     public static void main( String [] args )
7     {
8         // create shared object used by threads
9         SynchronizedBuffer sharedLocation = new SynchronizedBuffer();
10
11         // Display column heads for output
12         StringBuffer columnHeads =
13             new StringBuffer( "Operation" );
14         columnHeads.setLength( 40 );
15         columnHeads.append( "Buffer\t\tOccupied Count" );
16         System.err.println( columnHeads );
17         System.err.println();
18         sharedLocation.displayState( "Initial State" );
19
20         // create producer and consumer objects
21         Producer producer = new Producer( sharedLocation );
22         Consumer consumer = new Consumer( sharedLocation );
23
24         producer.start(); // start producer thread
```

6.4 Java Multithreading Case Study, Part III: Producer/Consumer Relationship in Java

```
25         consumer.start(); // start consumer thread
26
27     } // end main
28
29 } // end class SharedBufferTest2
```


6.4 Java Multithreading Case Study, Part III: Producer/Consumer Relationship in Java

Sample Output 1:

Operation	Buffer	Occupied Count
Initial State	-1	0
Consumer tries to read. Buffer empty. Consumer waits.	-1	0
Producer writes 1	1	1
Consumer reads 1	1	0
Consumer tries to read. Buffer empty. Consumer waits.	1	0
Producer writes 2	2	1
Consumer reads 2	2	0
Producer writes 3	3	1

6.4 Java Multithreading Case Study, Part III: Producer/Consumer Relationship in Java

Consumer reads 3	3	0
Consumer tries to read. Buffer empty. Consumer waits.	3	0
Producer writes 4	4	1
Consumer reads 4 Producer done producing. Terminating Producer.	4	0
Consumer read values totaling: 10. Terminating Consumer.		

6.4 Java Multithreading Case Study, Part III: Producer/Consumer Relationship in Java

Sample Output 2:

Operation	Buffer	Occupied Count
Initial State	-1	0
Consumer tries to read. Buffer empty. Consumer waits.	-1	0
Producer writes 1	1	1
Consumer reads 1	1	0
Producer writes 2	2	1
Producer tries to write. Buffer full. Producer waits.	2	1
Consumer reads 2	2	0
Producer writes 3	3	1

6.4 Java Multithreading Case Study, Part III: Producer/Consumer Relationship in Java

Consumer reads 3	3	0
Producer writes 4	4	1
Producer done producing. Terminating Producer.		
Consumer reads 4	4	0
Consumer read values totaling: 10. Terminating Consumer.		

Sample Output 3:

Operation	Buffer	Occupied Count
Initial State	-1	0
Producer writes 1	1	1

6.4 Java Multithreading Case Study, Part III: Producer/Consumer Relationship in Java

Sample Output 3 (Cont.):

Operation	Buffer	Occupied Count
Initial State	-1	0
Producer writes 1	1	1
Consumer reads 1	1	0
Producer writes 2	2	1
Consumer reads 2	2	0
Producer writes 3	3	1
Consumer reads 3	3	0
Producer writes 4	4	1
Producer done producing. Terminating Producer.		

6.4 Java Multithreading Case Study, Part III: Producer/Consumer Relationship in Java

Consumer reads 4

4

0

Consumer read values totaling: 10.
Terminating Consumer.